


# Deep Learning Topics: Special Connections, Encoder-Decoders, Transformers

Paul Rodriguez, PhD

(SDSC)

August 2024



# Outline

- **Basic word prediction task and motivating the attention strategy**
- **From Embeddings and Attention Head to Transformers**
- **BERT and GPT strategies**
- **Transformers in Science Applications**

# Dependences of Language

Consider this sequence:

*The Law will never be perfect, but it's application  
should be just - this is what we are missing, in my  
opinion <End of Sequence>*

What does 'it' refer to that can have an 'application'?

# Dependences of Language


Consider this sequence:

*The Law will never be perfect, but it's application should be just - this is what we are missing, in my opinion <End of Sequence>*

What does 'it' refer to that can have an 'application'?

e.g 'it' refers back to 'Law', which is part of 'the Law' noun phrase, which is the entity that will 'never be perfect', and so on ...

many dependencies and interactions



# A toy problem to get some intuition

- Let's use the following list of 5 tokens (ie words):  
    <**s**tart>, the, man, chicken, ordered
- Let's use this sequence of 6 tokens as our only sentence:  
    <start> the man ordered the chicken
- If we use **token** ids 1 to 5, the sequence is [1,2,3,5,2,4]
- **Now let's try to predict the next word by 'attention' idea**

# The toy task: predict next word

The data: 5 tokens ( $V=5$ ),

1 sequence (length= $T=6$ ): <Start> the man ordered the chicken

A basic solution is a bigram matrix:

**X=Sequence-to-Word, size is  $T \times V$**

Token  
Sequence

Next Token (ie word) Prediction

Pos	Word	<strt>	The	Man	Chikn	Order
0	<start>		1.0			
1	The			0.5	0.5	
2	Man					1.0
3	Orde.r		1.0			
4	The			0.5	0.5	
5	Chick.	1.0				



# The toy task: predict next word

The data: 5 tokens ( $V=5$ ),  
1 sequence (length= $T=6$ ):  $\langle \text{Start} \rangle$  the man ordered the chicken

A basic solution is a bigram matrix:  
**X=Sequence-to-Word, size is  $T \times V$**

*Challenge, can we learn predictions ( $\rightarrow$ )  
that depend on context of other tokens  
and/or position*

*After  $\langle \text{Start} \rangle$  the  $\rightarrow$  man = 1.0*

*After 'Ordered' the  $\rightarrow$  chicken = 1.0*

Token  
Sequence

Next Token (ie word) Prediction

Pos	Word	$\langle \text{str} \rangle$	The	Man	Chikn	Order
0	$\langle \text{start} \rangle$		1.0			
1	The			0.5	0.5	
2	Man					1.0
3	Order		1.0			
4	The			0.5	0.5	
5	Chick.	1.0				

# The attention idea

Let's get all tokens to 'pass information' about dependencies

E.G. for  $X$  a  $T \times V$  matrix of bigrams, we want to transform  $X$

**$X$  = Sequence-to-Word is  $T \times V$**

$$\begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 1.0 & 0 & 0 & 0 & 0 \end{pmatrix} \rightarrow \rightarrow \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 \end{pmatrix}$$



# The attention idea

Let's get all tokens to 'pass information' about dependencies

E.G. for  $X$  a  $T \times V$  matrix of bigrams, we want to transform  $X$   
we want  $W$  a  $T \times T$  matrix – aka 'attention' weights

$$\begin{array}{l} \text{W dependencies is } T \times T \\ \begin{pmatrix} w_{11} & \cdots & w_{1T} \\ \vdots & \vdots & \vdots \\ w_{T1} & \cdots & w_{TT} \end{pmatrix} \end{array} * \begin{array}{l} \text{X= Sequence-to-Word is } T \times V \\ \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 1.0 & 0 & 0 & 0 & 0 \end{pmatrix} \end{array} \rightarrow \rightarrow \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 \end{pmatrix}$$

# The attention idea

Let's get all tokens to 'pass information' about dependencies

E.G. for  $X$  a  $T \times V$  matrix of bigrams, we want to transform  $X$   
we want  $W$  a  $T \times T$  matrix – aka 'attention' weights

$$\begin{array}{l} \text{W dependencies is } T \times T \\ \begin{pmatrix} w_{11} & \cdots & w_{1T} \\ \vdots & \vdots & \vdots \\ w_{T1} & \cdots & w_{TT} \end{pmatrix} \end{array} * \begin{array}{l} \text{X= Sequence-to-Word is } T \times V \\ \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0.5 & 0.5 & 0 \\ 1.0 & 0 & 0 & 0 & 0 \end{pmatrix} \end{array} \rightarrow \rightarrow \begin{pmatrix} 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 & 0 \end{pmatrix}$$

*$W$  should reflect the interdependencies of the sequence.*


*Q: Where should  $W$  come from?*

- **Let's build up the attention architecture**

# An embedding layer for tokens

1. For an input vector  $X$  of  $V$  elements let the token-id determine which element is 1, the rest are 0.

2. Use a single, smaller, hidden layer with linear activation, i.e:  $Y = W * X$



# An embedding layer for tokens

1. For an input vector  $X$  of  $V$  elements let the token-id determine which element is 1, the rest are 0.


2. Use a single, smaller, hidden layer with linear activation, i.e:  $Y = W * X$

*Thus, each token id is converted to a lower dimensional vector with size  $1 \times E$*

*Each token sequence of vectors is  $T \times E$*



**X is  $V \times 1$**



# Get input embeddings

First, get sequence of token embeddings (call it X)

$$[1,2,3,5,2,4] \rightarrow X_{TxE}$$

# Get input and add in position info

First, get sequence of token embeddings (call it  $X$ )

$$[1,2,3,5,2,4] \rightarrow X_{TxE}$$

Then do the same for positions  $1 \dots T$

$$[1,2,3,4,5,6] \rightarrow P_{TxE}$$

*$X + P$  is final  $T \times E$  matrix of input embeddings*



# Embeddings for attention weights

Take Input Embeddings and build a 'Query' (Q) and 'Key' (K) embedding matrix of size  $T \times E$

$$\begin{aligned} X + P &\rightarrow Q_{T \times E} \\ X + P &\rightarrow K_{T \times E} \end{aligned}$$

Recall that embedding layers helps transform inputs into lower dimensional representations that capture information

# Embeddings for attention weights

Take Input Embeddings and build a 'Query' (Q) and 'Key' (K) embedding matrix of size  $T \times E$

$$\begin{aligned} X + P &\rightarrow Q_{T \times E} \\ X + P &\rightarrow K_{T \times E} \end{aligned}$$

Now let  $W = Q * K'$

Recall that embedding layers helps transform inputs into lower dimensional representations that capture information

Notice that every token's embedding gets to 'interact' with every other token's embedding to make up the  $T \times T$  elements of  $W$

# Embeddings for attention weights

Take Input Embeddings and build a 'Query' (Q) and 'Key' (K) embedding matrix of size  $T \times E$

$$\begin{aligned} X + P &\rightarrow Q_{T \times E} \\ X + P &\rightarrow K_{T \times E} \end{aligned}$$

Recall that embedding layers helps transform inputs into lower dimensional representations that capture information

Now let  $W = Q * K'$

Notice that every token's embedding gets to 'interact' with every other token's embedding to make up the  $T \times T$  elements of  $W$

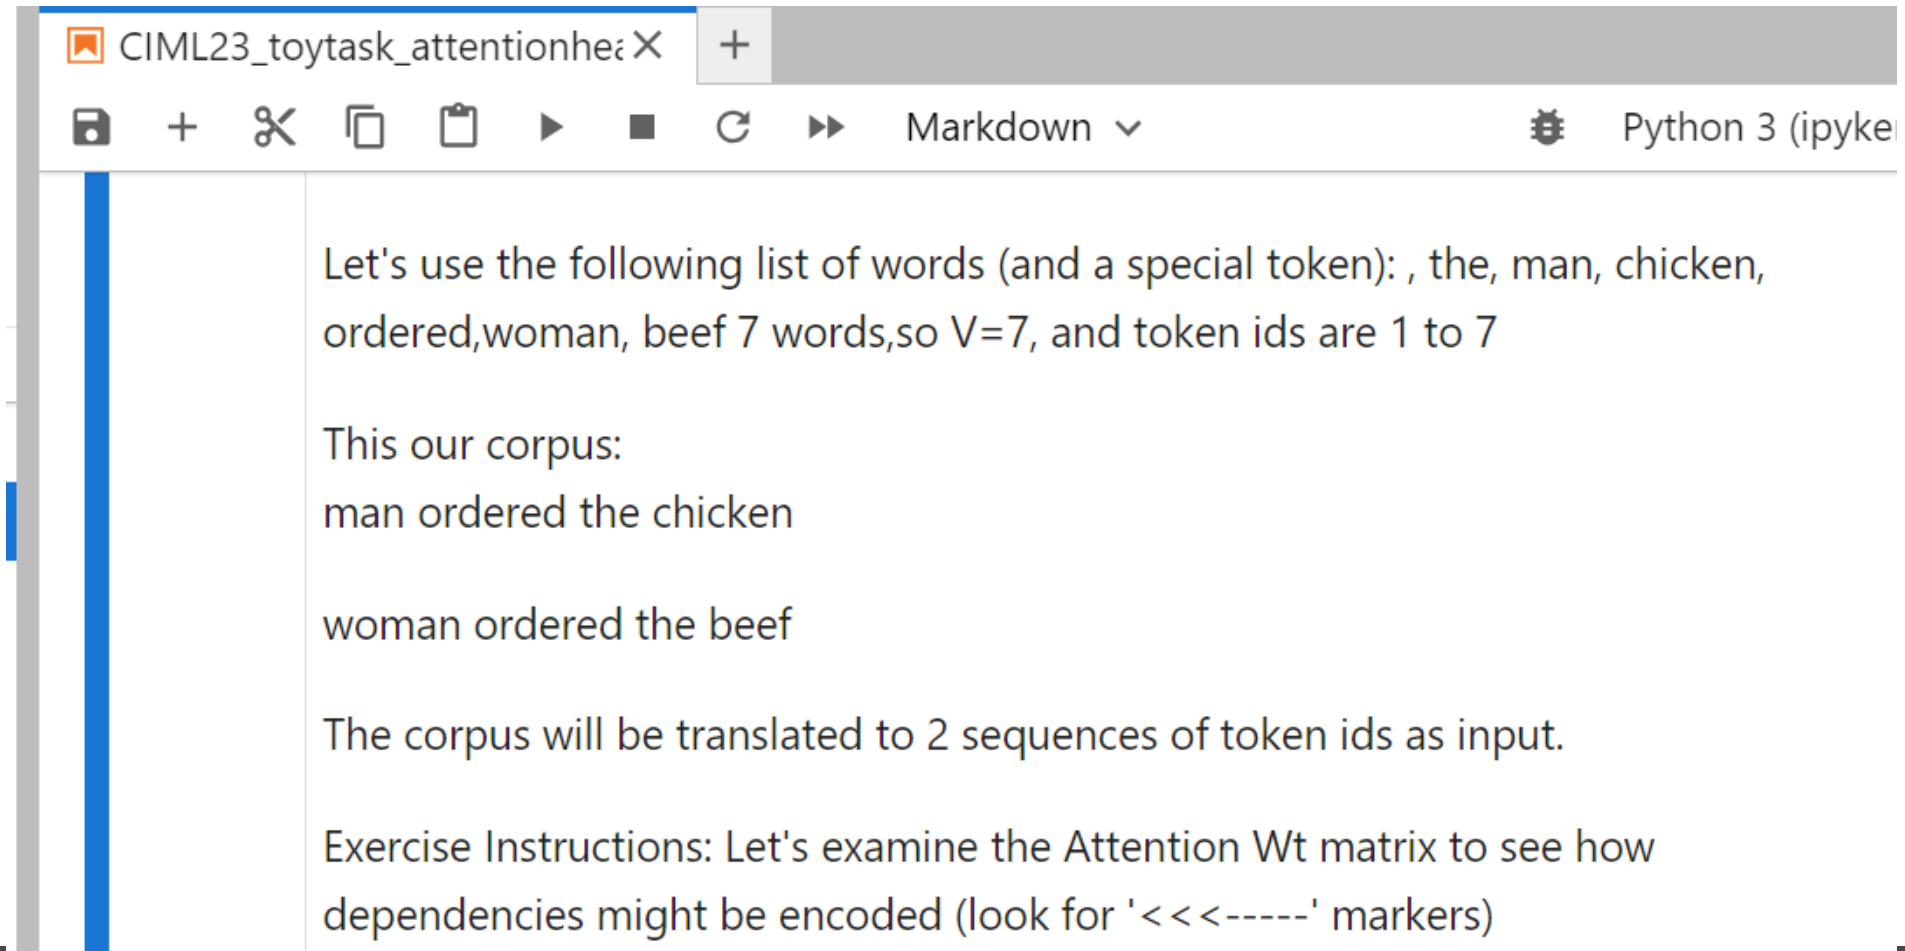
$$X + P \rightarrow V_{T \times E}$$

Finally, instead of a pre-built bigram matrix, use another embedding for a 'Value' V matrix

Now, we can take  $W * V$

An example of attention head with a toy task:

Run the “toytask\_attention notebook” and observe the printed predictions and attention weights. Try changing H – does it help/hurt?



The screenshot shows a Jupyter Notebook window titled "CIML23\_toytask\_attentionhe". The toolbar includes icons for saving, adding, deleting, copying, pasting, running, and a dropdown menu for "Markdown". The notebook content is as follows:

```
Let's use the following list of words (and a special token): , the, man, chicken,
ordered,woman, beef 7 words,so V=7, and token ids are 1 to 7

This our corpus:
man ordered the chicken

woman ordered the beef

The corpus will be translated to 2 sequences of token ids as input.

Exercise Instructions: Let's examine the Attention Wt matrix to see how
dependencies might be encoded (look for '<<<-----' markers)
```

## Output TxV predictions:

Notice that the → [chicken or beef] predictions change depending on who's ordering

	<ST>	the	man	chkn	ordrd	woman	beef
0 <ST>	0.02	0.18	0.98	0.38	0.16	0.98	0.04
1 man	0.30	0.22	0.16	0.62	0.99	0.12	0.35
2 ordrd	0.68	1.00	0.54	0.37	0.51	0.17	0.58
3 the	0.25	0.07	0.05	0.99	0.58	0.11	0.85
4 chkn	1.00	0.51	0.09	0.23	0.51	0.08	0.56

	<ST>	the	man	chkn	ordrd	woman	beef
0 <ST>	0.02	0.18	0.98	0.38	0.16	0.98	0.04
1 woman	0.42	0.23	0.01	0.07	1.00	0.17	0.85
2 ordrd	0.68	1.00	0.54	0.36	0.51	0.17	0.58
3 the	0.22	0.09	0.01	0.80	0.74	0.03	0.99
4 beef	1.00	0.52	0.08	0.22	0.49	0.07	0.59

## Output TxV predictions:

Notice that the → [chicken or beef] predictions change depending on who's ordering

	<ST>	the	man	chkn	ordrd	woman	beef
0 <ST>	0.02	0.18	0.98	0.38	0.16	0.98	0.04
1 man	0.30	0.22	0.16	0.62	0.99	0.12	0.35
2 ordrd	0.68	1.00	0.54	0.37	0.51	0.17	0.58
3 the	0.25	0.07	0.05	0.99	0.58	0.11	0.85
4 chkn	1.00	0.51	0.09	0.23	0.51	0.08	0.56

	<ST>	the	man	chkn	ordrd	woman	beef
0 <ST>	0.02	0.18	0.98	0.38	0.16	0.98	0.04
1 woman	0.42	0.23	0.01	0.07	1.00	0.17	0.85
2 ordrd	0.68	1.00	0.54	0.36	0.51	0.17	0.58
3 the	0.22	0.09	0.01	0.80	0.74	0.03	0.99
4 beef	1.00	0.52	0.08	0.22	0.49	0.07	0.59

## TxT Attn Wts:

Notice weights for “the” are different for “man” or “woman”

	0 <ST>	1 man	2 ordrd	3 the	4 chkn
0 <ST>	1.00	0.00	0.00	0.00	0.00
1 man	0.06	0.94	0.00	0.00	0.00
2 ordrd	0.06	0.00	0.93	0.00	0.00
3 the	0.01	0.78	0.01	0.20	0.00
4 chkn	0.02	0.01	0.04	0.04	0.90

	0 <ST>	1 woman	2 ordrd	3 the	4 beef
0 <ST>	1.00	0.00	0.00	0.00	0.00
1 woman	0.10	0.90	0.00	0.00	0.00
2 ordrd	0.06	0.00	0.93	0.00	0.00
3 the	0.01	0.87	0.01	0.12	0.00
4 beef	0.01	0.02	0.04	0.04	0.89

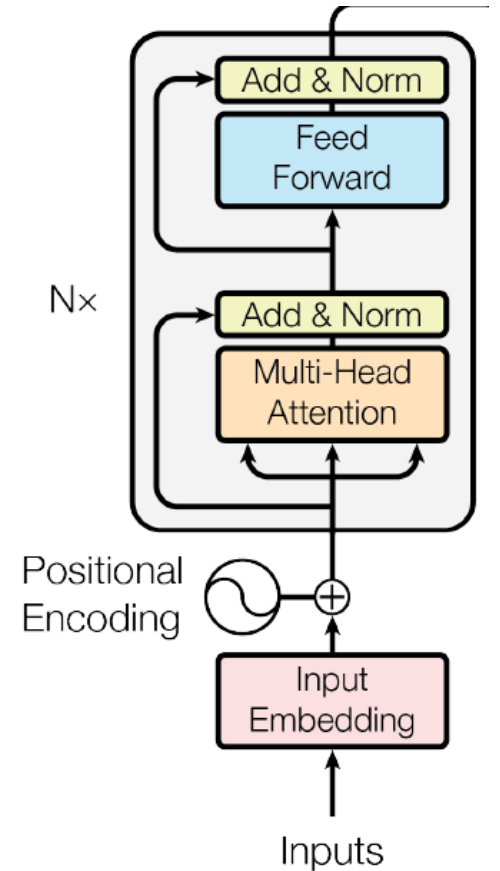
- **pause**



# Finally, a transformer

Include skip-add connections

Include Layer Normalization or DropOut layers

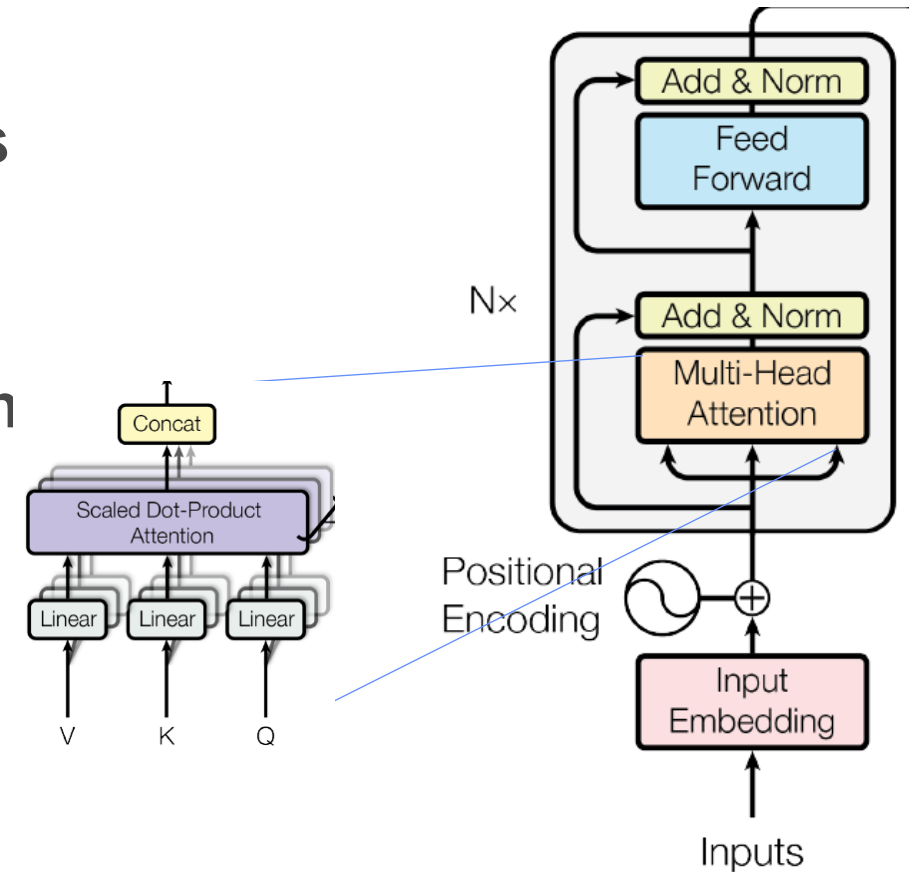


# Finally, a transformer

Include skip-add connections

Include Layer Normalization or DropOut layers

Multi-Head – for N heads produce  $T \times (E/N)$  each



# Finally, a transformer


Include skip-add connections

Include Layer Normalization or DropOut layers

Multi-Head – for  $N$  heads produce  $Tx(E/N)$  each

Add MLP layers on top and keep output  $TxE$

***Then stack transformers!***



# Finally, a transformer

Include skip-add connections

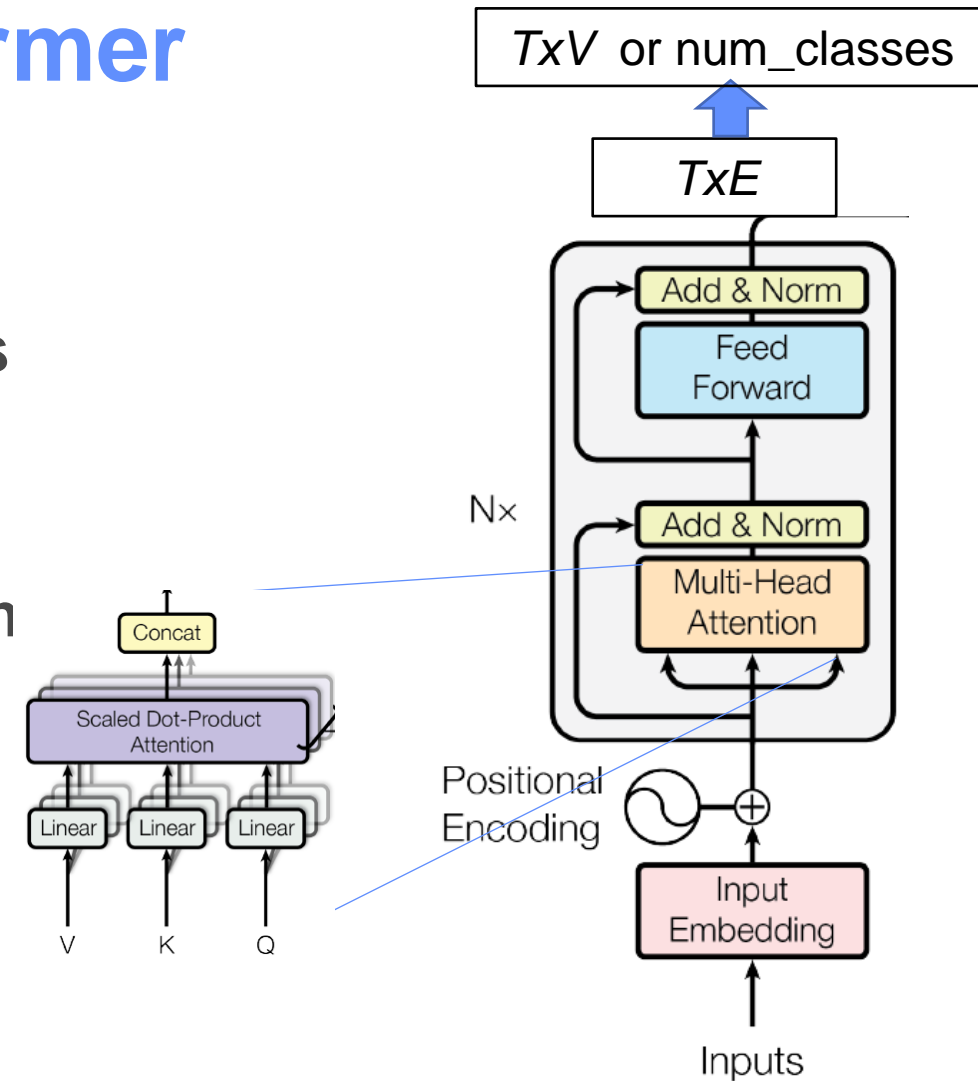
Include Layer Normalization or DropOut layers

Multi-Head – for  $N$  heads produce  $Tx(E/N)$  each

Add MLP layers on top and keep output  $TxE$

***Then stack transformers!***

Finally, produce some classification or word predictions




## 2 kinds of training strategies

**GPT – predict next word only look back at prior context (which could be a whole document)**

*Put mask on attention weights so that predictions only depend on previous tokens*

**BERT – *No attention mask* so all token dependencies can influence all other tokens predictions**

**Special tokens help create a variety of tasks**

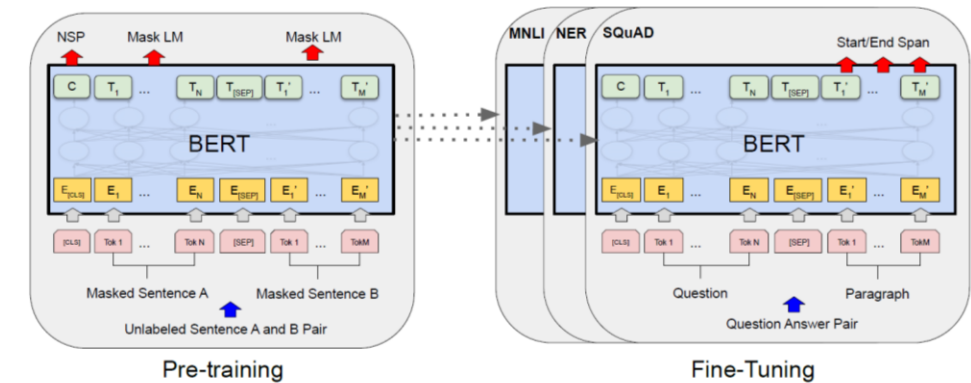


# BERT (Bidirectional Encoder Representations from Transformers)

Goal: Train a model to develop general token-level AND sentence-level encoding

1 Pretrain on:

- fill-in-the-blank
- binary classification if 2 sentences go together



Devlin, etal, 2019

2 Fine tune on variety of tasks

# GPT (generative pre-trained transformer)

Goal: Train a transformer model at large scale so that it develops very general representations that are useful for many language tasks.

‘GPT3 shows strong performance on many NLP tasks so that with a few examples it nearly matches fine-tuned systems’

Lang. Models are Few Shot Learners  
Brown, et al, 2020, openAI,

‘Instruction tuning’ improves models so they don’t need examples

Finetuned language models are zero-shot Learners. Wei et al, 2022, Google

## Finetune on many tasks (“instruction-tuning”)

### Input (Commonsense Reasoning)

Here is a goal: Get a cool sleep on summer days.  
How would you accomplish this goal?  
OPTIONS:  
-Keep stack of pillow cases in fridge.  
-Keep stack of pillow cases in oven.

### Target

keep stack of pillow cases in fridge

### Input (Translation)

Translate this sentence to Spanish:  
The new office building was built in less than three months.

### Target

El nuevo edificio de oficinas se construyó en tres meses.

Sentiment analysis tasks

Coreference resolution tasks

...



- **Transformers for Science applications**

**Can anything be cast as a kind of sentence, or an arrangement of tokens?**


# ESM Fold model

*Language models of protein sequences at the scale of evolution enable accurate structure prediction*

Lin et al, Meta Research 2022

Atom level structure prediction

Uses protein sequence as input to transformer layers (like LLM)



# ESM Fold model

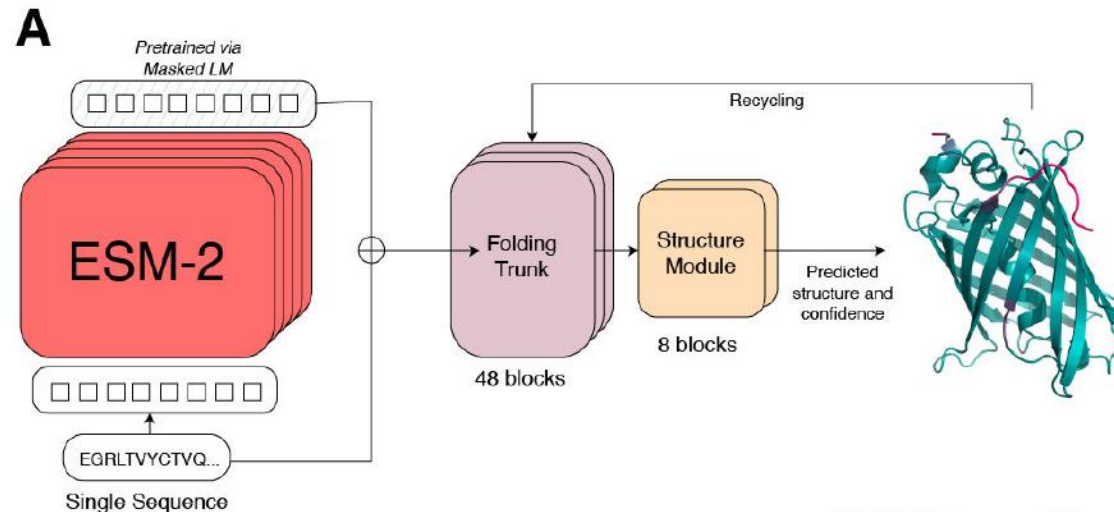
*Language models of protein sequences at the scale of evolution enable accurate structure prediction*

Lin et al, Meta Research 2022

## Atom level structure prediction

Uses protein sequence as input to transformer layers (like LLM)

Predicts a map of protein contact which gets *iteratively refined* by a 'folding block' transformer and structure module (similar to AlphaFold2, but faster)




# Molecule Attention Transformer

(Maziarka et al. 2020)

Molecular property prediction

Uses the set of atoms as input (like sentence tokens)

Includes spatial information by using a sum of the attention matrix, a distance matrix, and an adjacency matrix.



# The Visual Transformer (ViT)

*An image is worth 16x16 words:*

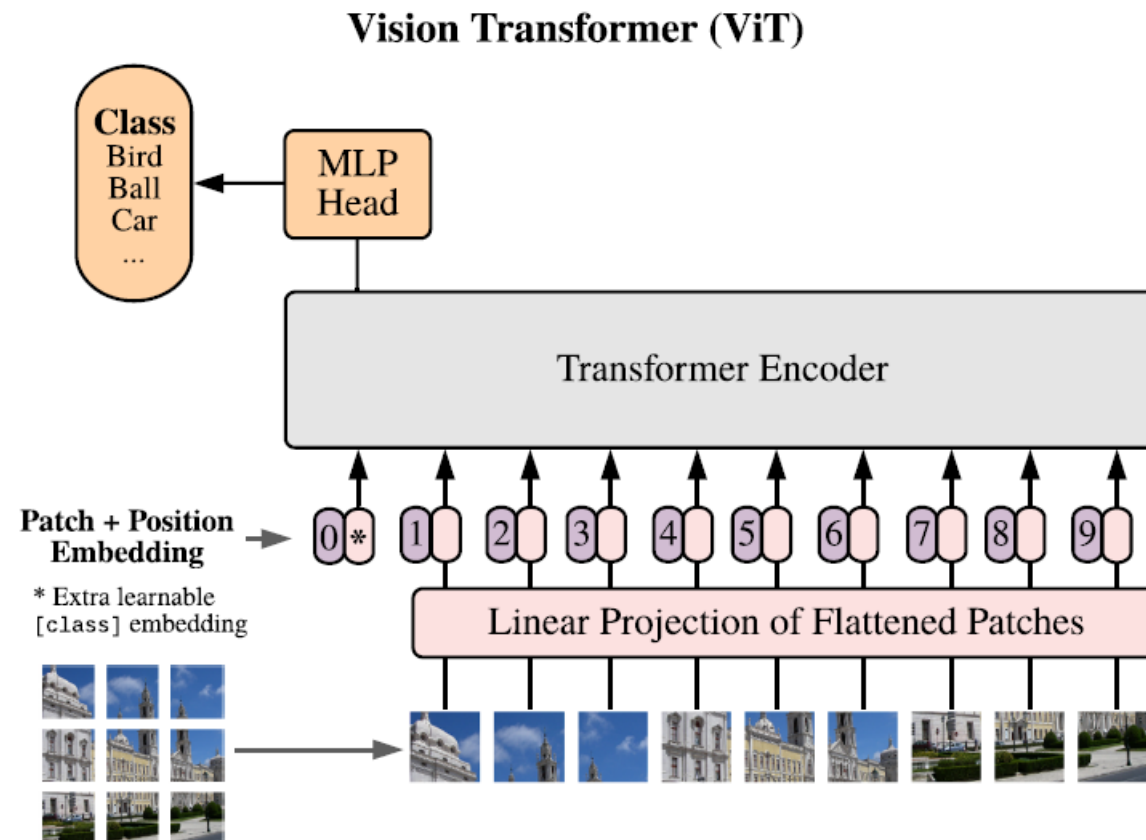
*Transformers for image recognition at scale*

Adosovitski, et al, 2021, Google Research

Uses a sequences of image patches (16x16)  
like a sentence of tokens (ie 224x224 pixels is  
16x16 patches of 14x14 pixels)

Uses a classification token like Bert to learn  
image output classes

Competitive or better than CNNs but might need  
more data



- **Combining images and text often makes DL work better, or more generic, for image or text tasks**

# CLIP – Contrastive Language-Image Pretraining

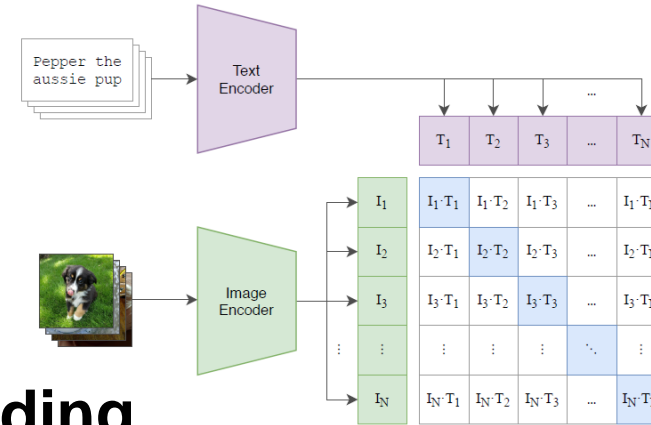
## *Learning Transferable Visual Models From Natural Language Supervision*

Radford et al, 2021, Open AI

Uses 400M images and captions for training

Learns a combination of text and image embeddings into a new **multi-modal embedding**

(1) Contrastive pre-training





# CLIP – Contrastive Language-Image Pretraining

## *Learning Transferable Visual Models From Natural Language Supervision*

Radford et al, 2021, Open AI


Uses 400M images and captions for training

Learns a combination of text and image embeddings into a new **multi-modal embedding**

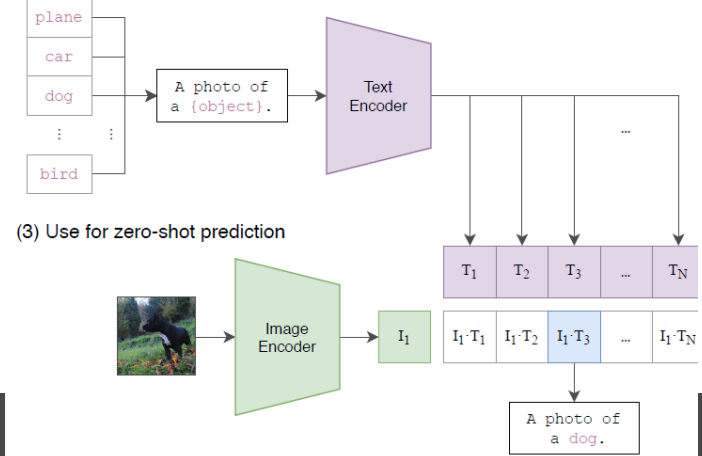
Performs classification by prompting it with an image and possible captions

Note: CLIP with diffusions gets close to DALL-E

(1) Contrastive pre-training






(2) Create dataset classifier from label text



# BIOCLIP: A Vision Foundation Model for the Tree of Life

Stevens, etal 2024 OSU

- Uses pre-trained CLIP for a base
- Uses Tree-of-Life 10M dataset of biology images with taxonomic labels
- The taxonomic hierarchy is presented as a sequence of words for different species



end